

MPHYG002: Research Computing with C++

Parallel Implementation of the Game of Life

Michael Ebner
(Student Number: 1500 1660)

April 26, 2015

Contents

1	Conway’s Game of Life	1
2	Implementation	2
2.1	Structure of project-folder	2
2.2	Choice of design	2
2.3	Parallelisation	2
2.4	Unit tests	4
2.5	Build, run and test	4
2.5.1	Build instructions	5
2.5.2	Running the programme	5
2.5.3	Run unit tests	5
2.6	Results	5
2.6.1	Video	6
2.6.2	Parallel vs. serial implementation	6

1 Conway's Game of Life

Based on the lecture notes and http://en.wikipedia.org/wiki/Conways_Game_of_Life¹ a summary of the rules of Conway's Game of Life reads as follows:

- The game is represented by an infinite two-dimensional, rectangular grid of cells.
- Every cell can represent two possible states: dead or alive.
- The grid with its cells represent a population at a discrete point of time.
- The state of each cell in the consecutive time step depends on its 8 adjacent neighbours.
- The transition of the each cell's state are described by these rules, cf. table 1.1:
 - A live cell remains alive in case it is surrounded by either two or three living cells. Otherwise it dies which relates to the cases of under-population or overcrowding, respectively.
 - A dead cell with exactly three surrounding living cells becomes alive associated with the idea of reproduction.

		Number of neighbour cells alive								
		0	1	2	3	4	5	6	7	8
Current state	0	0	0	0	1	0	0	0	0	0
	1	0	0	1	1	0	0	0	0	0

Table 1.1: Transition rules whereby "0" indicates a dead and "1" a live cell

¹Retrieved: April 25, October 2015

2 Implementation

2.1 Structure of project-folder

The project can be downloaded from <https://github.com/renbem/RCCPP-coursework02>. The code is structured in several folders within the main directory:

- **documentation:** Contains all documentation of the project including this report and the code documentation generated by doxygen¹.
- **include:** Contains all header files (`*.h`)
- **matlab:** Contains the MATLAB-files used for the random generation of initial boards (or grid according to the terminology used in chapter 1), the creation of the Game of Life video² and statistical evaluation of the results.
- **source:** Contains the code files (`*.cc` including `main.cc`).
- **test:** Contains all data and the source file `main_UnitTests.cc` for the unit tests.

2.2 Choice of design

The implementation follows the design as illustrated in fig. 2.1 which translates the concept of a having a game consisting of boards which in turn consist of cells.

The classes are designed so that a new iteration is triggered by the member function `computeNextStep` of `Game`, cf. fig. 2.2. I decided to connect the "rules of the game" with the board via the two methods `determineNeighbourCells` and `applyTransitionRules` since they constitute two separate logical units. However, one could argue that this might be better situated in the class `Game` since the rules can be considered as an intrinsic property of the game and the board could be of arbitrary shape³. Nevertheless, since Conway's Game of Life relies on well-balanced rules in order to have a "meaningful game" (i.e. no extinction or overcrowding of cells) the definition of a neighbour and the propagation rules depend strongly on the shape and therefore on the board.

2.3 Parallelisation

The code was parallelised via **OpenMP** within the method `computeNextStep`, cf. line 10 in listing 2.1. The line 9 and 11 was added since I could not compile it without any error messages on Mac (further information in section 2.5) but it works on Ubuntu. Therefore, the code runs on non-linux platforms only in the serial version.

¹www.doxygen.org

²Provided at https://www.youtube.com/watch?v=AeIm2I_9n5g&feature=youtu.be

³E.g. circular, hexagonal, x -dimensional, etc.

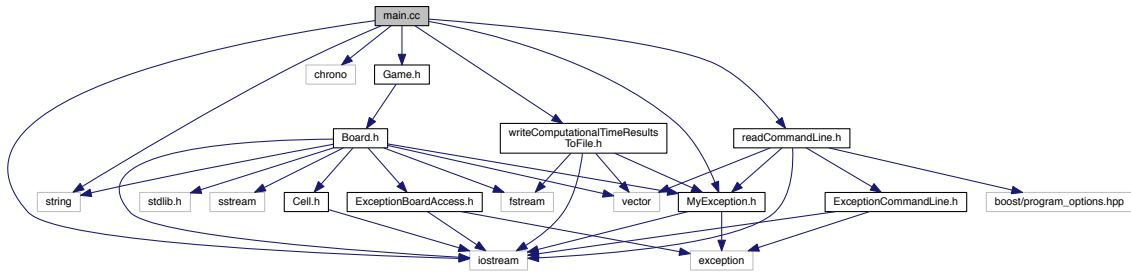


Figure 2.1: Overview of implementation

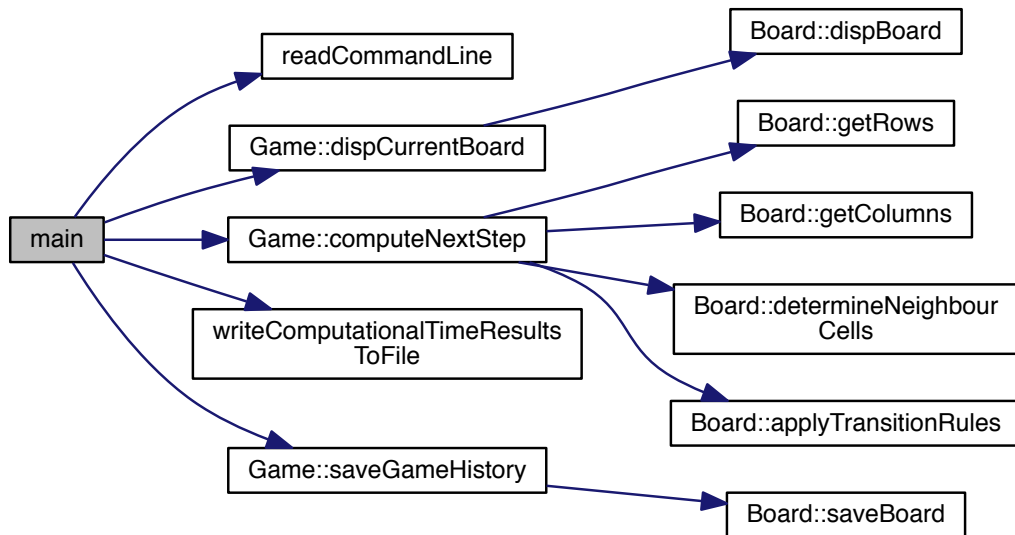


Figure 2.2: Call graph of `main.cc`

Listing 2.1: Game::computeNextStep

```
1 void Game::computeNextStep(){
2     Board currentBoard = boardHistory.back();
3     Board updatedBoard = currentBoard;
4
5     unsigned int irows = currentBoard.getRows();
6     unsigned int icolumns = currentBoard.getColumns();
7     std::vector<Cell> neighbourCells;
8
9     #ifdef __linux
10        #pragma omp parallel for private(neighbourCells)
11    #endif
12    for (unsigned int row = 0; row < irows; ++row){
13
14        for (unsigned int col = 0; col < icolumns; ++col){
15            neighbourCells
16                = currentBoard.determineNeighbourCells(row, col);
17            updatedBoard.applyTransitionRules(row,col,neighbourCells);
18        }
19    }
20    this->pushBoard(updatedBoard);
21 }
```

2.4 Unit tests

Unit tests were written to test the following cases:

- Check input file⁴.
- Several tests to check malformed command lines.
- Several tests to check whether the correct neighbours are returned by `Board::neighbourCells`. (The "infinite board" was implemented by assuming periodic boundary conditions.)
- Check that attempted access to non-existing indices of board throws an exception.
- Test whether one step of Conway's Game of Life is computed correctly.
- Test whether the parallel computation via OpenMP returns the same board as the serial version.

The corresponding file `main_UnitTests.cc` is located in `test/`.

2.5 Build, run and test

The `CMakeLists.txt` file in the parent directory was tested on two operating systems:

- OS X 10.10.3 (Macbook Pro, 2.5 GHz Intel Core i7)

⁴It is only tested whether the input file exists. Many other cases could be tested in addition. E.g.: Format of file appropriate?, Number of columns always identical?, Minimum number of rows and columns alright? etc.

- Ubuntu 14.04 (Virtual machine via VirtualBox⁵ within OS X 10.10.3)

Unfortunately, I couldn't run OpenMP on Mac without any errors (comments are provided in the `CMakeLists.txt` file). Therefore, I used VirtualBox to build it on Ubuntu where everything went flawlessly. Consequently, I configured the CMake file (and added respective lines in `Game.cc`, e.g. lines 9 and 11 in listing 2.1) so that the code runs on Mac in serial and on Ubuntu with OpenMP. The results in section 2.6 are based on the computation in Ubuntu.

2.5.1 Build instructions

In order to build the code run the following lines in the main directory:

```
mkdir build
cd build
cmake ..
make
```

Two binary files will be generated:

- `conwaysGameOfLife` located in `build/bin/`
- `conwaysGameOfLife_UnitTests` located in `build/test/bin/`

2.5.2 Running the programme

For running the code several example data sets are provided in `build/test/exampleData/`. A possible command of the programme within the `build`-folder reads

```
bin/conwaysGameOfLife --i "test/exampleData/InitialBoardRandom_10Times10.txt"
--o "GameHistory.txt" --s 100
```

which uses the initial board defined in `InitialBoardRandom_10Times10.txt` and performs 100 steps by applying the transition rules given in table 1.1. A concatenated history of all computed boards is then stored in `GameHistory.txt` and the information of the corresponding computational time per iteration is saved in `GameHistory_ComputationalTime.txt`.

Alternatively any other file `InitialBoardRandom_*.txt` in `test/exampleData` can be used as input file⁶.

2.5.3 Run unit tests

To run the unit tests, execute the command

```
./conwaysGameOfLife_UnitTests
```

in `build/test/bin/`.

2.6 Results

In this section the results, obtained on Ubuntu, are shown and shall be discussed.

⁵<https://www.virtualbox.org/>

⁶A "on-screen-simulation" and its corresponding computational time can also be enabled by setting the respective flags `flagDisplayGame` and `flagDisplayComputationalTime` to `true` within `main.cc`.

```

→ time bin/conwaysGameOfLife --i "test/exampleData/InitialBoardRandom_200Times400.txt" -
-o "out_BoardHistory.txt" --s 1000
-----
Input file given (test/exampleData/InitialBoardRandom_200Times400.txt).
Output file for game history given (out_BoardHistory.txt).
Number of steps of game given (1000).
-----
Total time = 119.527451 s
bin/conwaysGameOfLife --i --o "out_BoardHistory.txt" --s 1000 120.77s user 8.73s system
97% cpu 2:12.52 total

```

(a) Serial implementation run

```

→ time bin/conwaysGameOfLife --i "test/exampleData/InitialBoardRandom_200Times400.txt" -
-o "out_BoardHistory.txt" --s 1000
-----
Input file given (test/exampleData/InitialBoardRandom_200Times400.txt).
Output file for game history given (out_BoardHistory.txt).
Number of steps of game given (1000).
-----
Total time = 41.201670 s
bin/conwaysGameOfLife --i --o "out_BoardHistory.txt" --s 1000 144.80s user 9.33s system
278% cpu 55.428 total

```

(b) Parallel implementation run

Figure 2.3: Statistic obtained via command 'time'

2.6.1 Video

An example board of dimension 200×400 was computed and uploaded to https://www.youtube.com/watch?v=AeIm2I_9n5g&feature=youtu.be. Several phenomena can be observed such as gliders, oscillators and stationary patterns.

2.6.2 Parallel vs. serial implementation

In this comparison the Game of Life rules were applied to boards of dimensions

- | | |
|-----------------------------------|--|
| (i) 10×10 (100 Cells) | (vi) 100×100 (10 000 Cells) |
| (ii) 10×20 (200 Cells) | (vii) 100×200 (20 000 Cells) |
| (iii) 10×50 (500 Cells) | (viii) 200×200 (40 000 Cells) |
| (iv) 10×100 (1000 Cells) | (ix) 200×300 (60 000 Cells) |
| (v) 50×100 (5000 Cells) | (x) 200×400 (80 000 Cells) |

for in total 1000 iterations for both the serial and parallel version. In fig. 2.3 the execution command and the corresponding results within the terminal are shown. The serial implementation (line 10 in listing 2.1 was commented) yielded the results shown in fig. 2.3a. The mere computation time (neglecting reading and writing processes) of performing 1000 steps on a 200×400 board took about 120 s. The whole execution (including reading and writing procedures) lasted, according to the `time` command, about 2:12 min. The serial computation also becomes apparent by noting the CPU load of 97%. The parallel results, shown in fig. 2.3b, illustrate the speed-up by using OpenMP. Time was reduced by a third and CPU load increased to 278%.

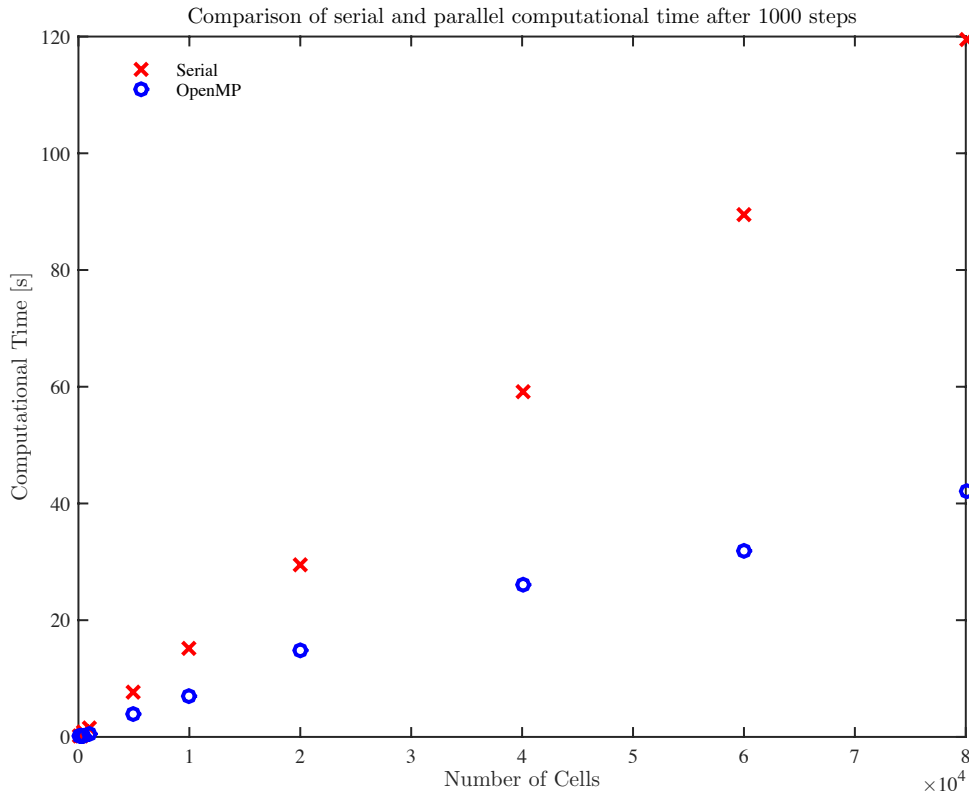


Figure 2.4: Computational time over number of cells

The respective cumulative computational times over the number of cells are illustrated in fig. 2.4 which states the expected roughly linear increase for each computational method. It can be observed that the implementation via OpenMP is already beneficial in lower numbers and becomes clearly evident in higher ones. However, the relative speed up seems to fluctuate.

In fig. 2.5 the computational time is shown over iterations for selected boards. Again, the higher the number of cells the more severe the impact of OpenMP on cumulative time.

Lastly, fig. 2.6 depicts the relative speed-up. In general, the advantage of using OpenMP becomes more apparent in higher number of cells. Since it was computed on 4 cores the theoretical maximum speed-up would be less than 4. However, the relative figures vary substantially over the number of cells. For a closer investigation larger boards could be investigated. Moreover, the fact that a virtual machine needed to be used could have had an impact on the results as well.

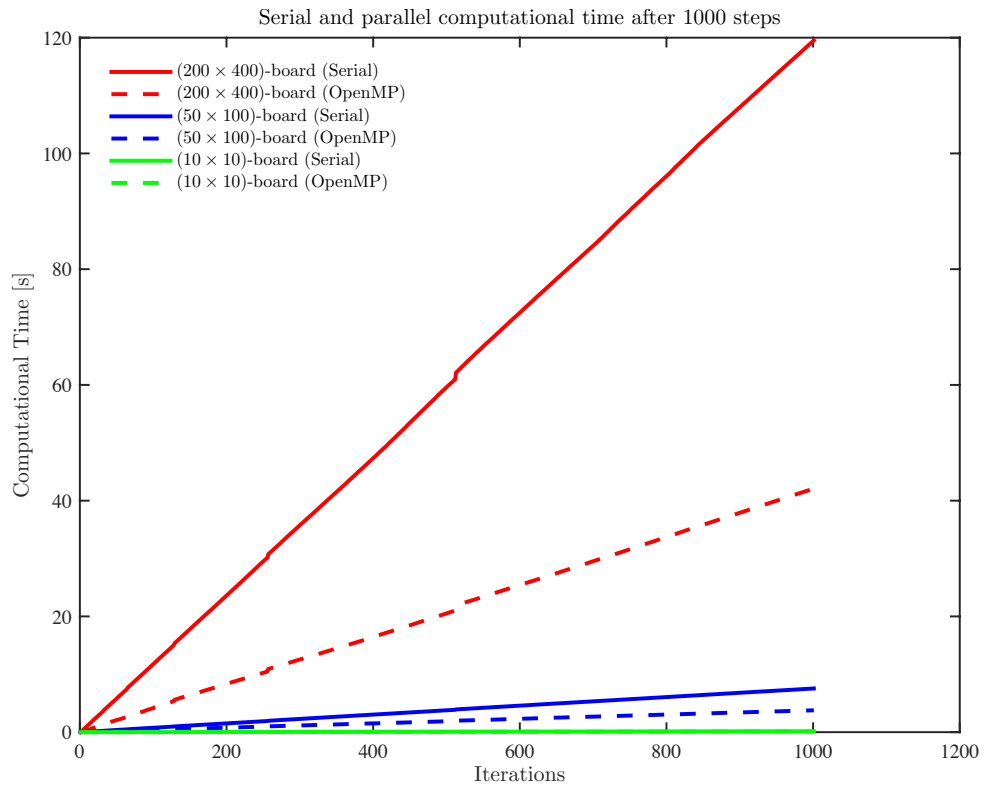


Figure 2.5: Computational time over number of iterations




Figure 2.6: Relative speed-up